
Week 10

**Project 3: An Introduction
to File Systems**

Classes COP4610 / CGS5765
Florida State University

Introduction

- The goal of project 3 is to understand
 - basic file system design and implementation
 - file system testing
 - data serialization/de-serialization
- At the end of the project, you will feel like a file system expert!

Outline

- Background
 - Setting up your environment
 - Mounting file systems
- Project 3
 - Specification
 - Downloading and testing file system image
 - General FAT32 data structures
 - Endian-ness

Environment Setup

Get ready for Project 3!

Project 3 Environment

- Must develop inside Linux environment with root access
 - Use Debian virtual machine or lab machine from project 2
- I will be grading your projects inside of my copy of the class virtual machine

Project Environment

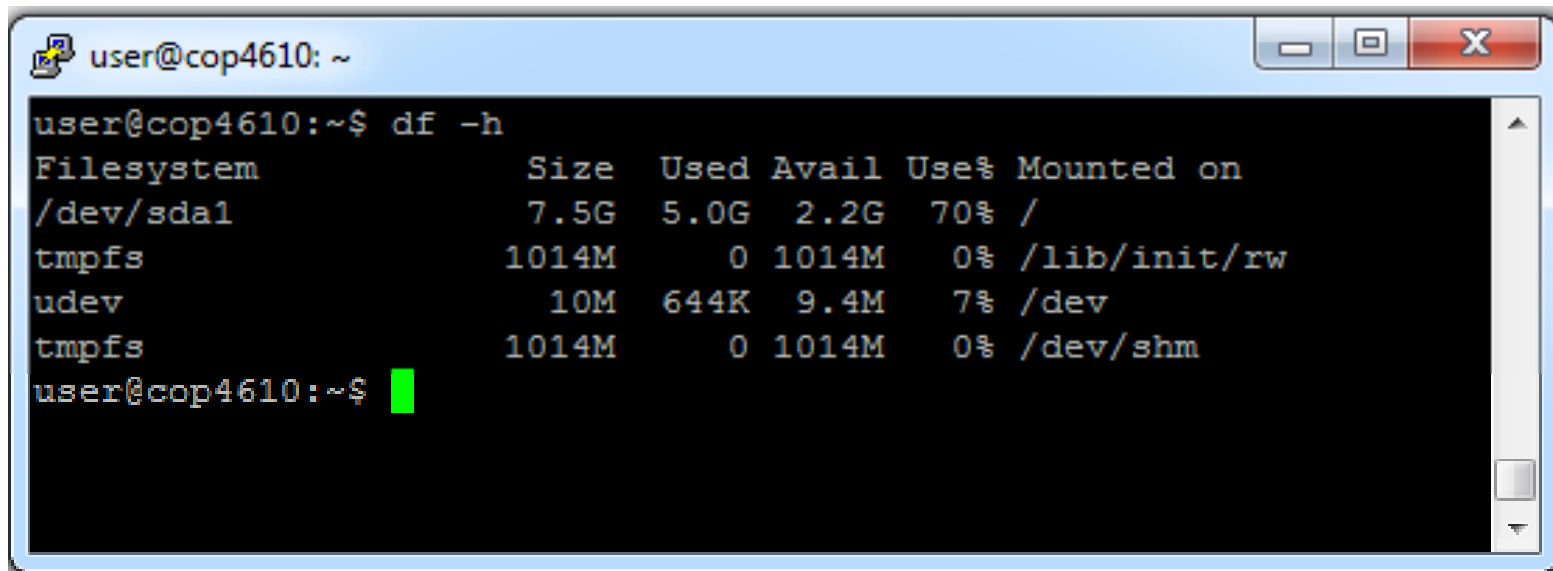
- Kernel version no longer matters
 - Can go back to the default 2.6.26 version if you trust it more
- Entire project will be in userspace, not kernel!
 - Please use debuggers, they will save you time
 - gdb, ddd, others...
- Programming language is still C

Running out of room?

- You must have at least 64MB free plus room for your source code
- To see how much room you have left inside your machine, issue the following command:

```
$> df -h
```

df -h



A terminal window titled 'user@cop4610: ~' showing the output of the 'df -h' command. The output is a table with columns: Filesystem, Size, Used, Avail, Use%, and Mounted on. The data shows that /dev/sda1 is the root filesystem with 7.5G total size, 5.0G used, and 2.2G available (70% used). Other filesystems like tmpfs and udev are also listed with their respective sizes and usage.

```
user@cop4610:~$ df -h
Filesystem      Size  Used Avail Use% Mounted on
/dev/sda1       7.5G  5.0G  2.2G  70% /
tmpfs           1014M    0 1014M   0% /lib/init/rw
udev            10M   644K   9.4M   7% /dev
tmpfs           1014M    0 1014M   0% /dev/shm
user@cop4610:~$
```

- /dev/sda is root file system mounted on “/”
 - Has 2.2GB currently available

Cleaning up space

- If you need to clean up some more space
 - Issue 'make clean' inside of your 2.6.32 kernel source
 - Delete the 2.6.26 kernel source

Mounting File Systems

Unix File Hierarchy

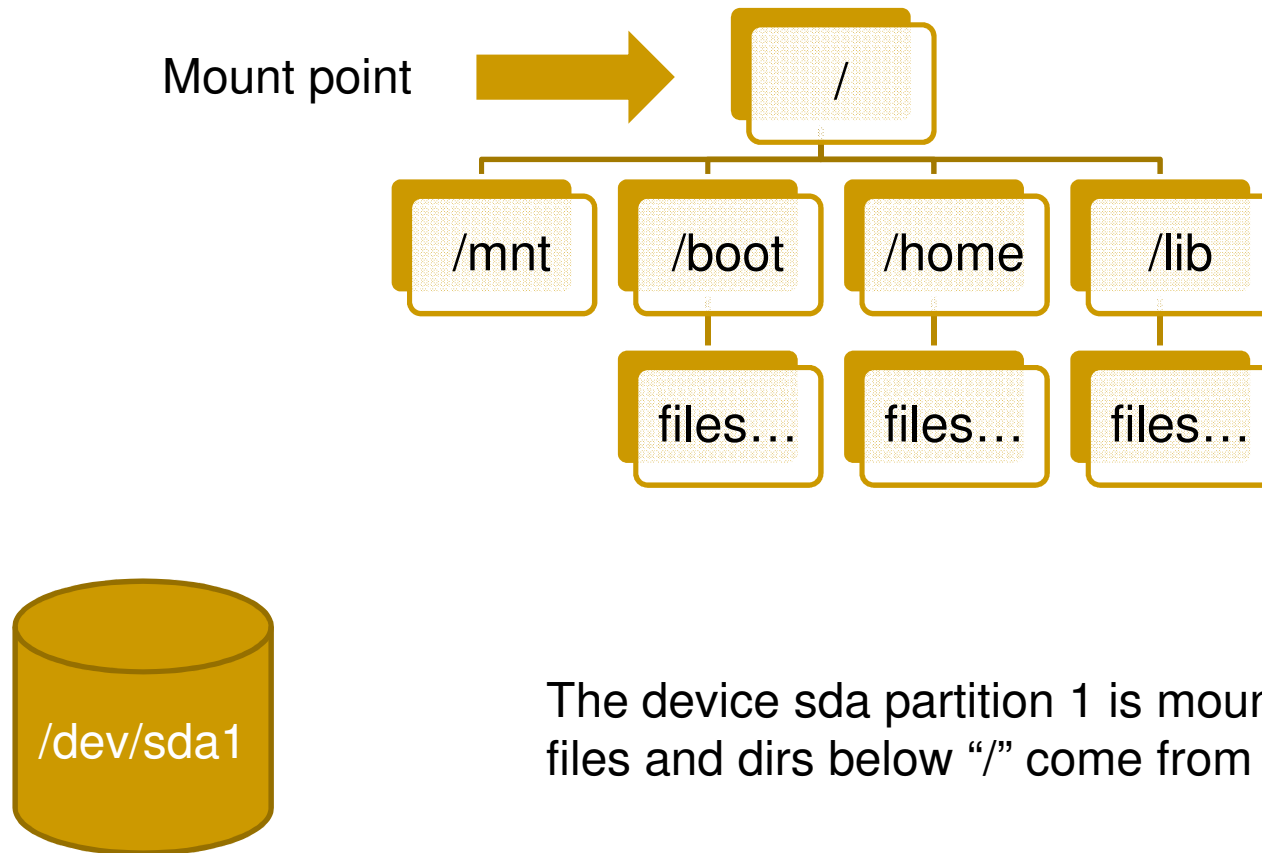
- All files accessible in a Unix system are arranged in one big tree
 - Also called the ***file hierarchy***
 - Tree is rooted (starts) at /
- These files can be spread out over several devices
- The **mount** command serves to attach the file system found on some device to the big file tree

'mount' command

```
■ mount  
■ mount <device> <mount directory>
```

- Typing 'mount' without arguments shows you what is mounted and where
- Second example attaches a device or partition to a directory
 - Must have root privileges

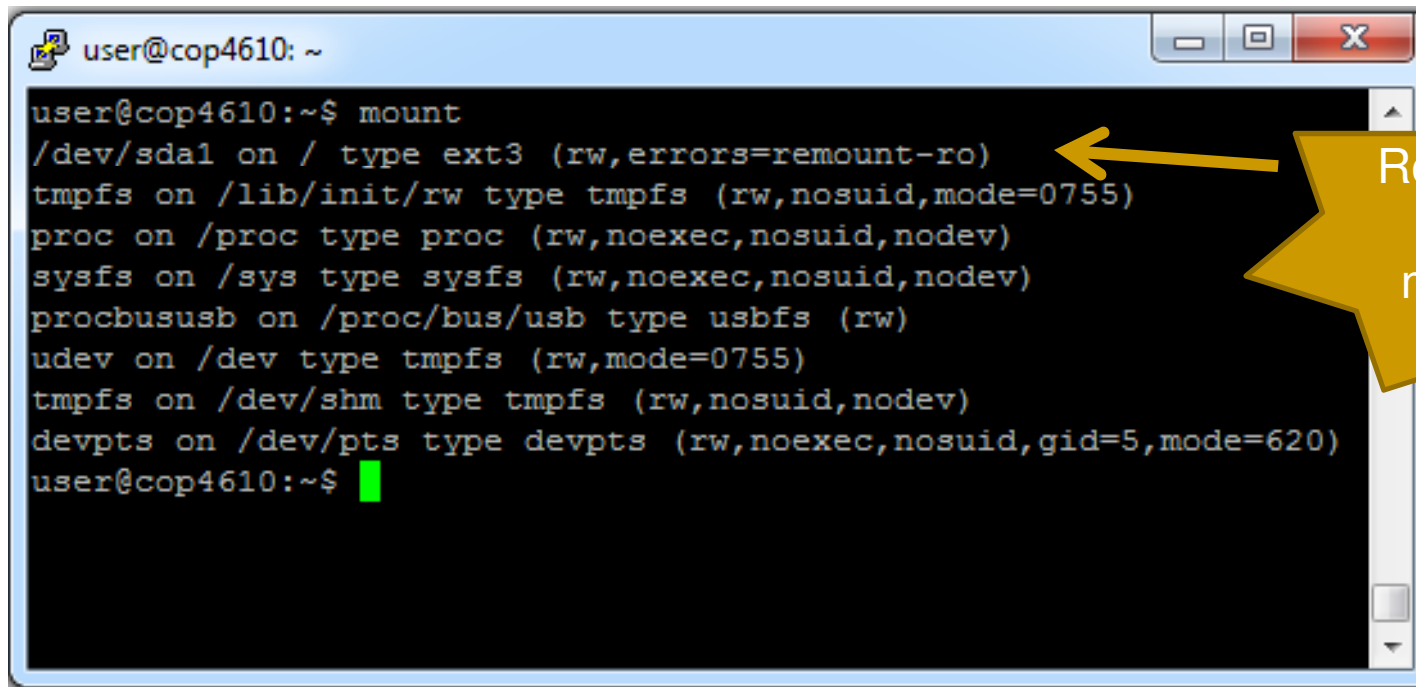
Mount Example



The device sda partition 1 is mounted at “/”. All files and dirs below “/” come from this device.

Mount Example



- Type command 'mount' without any arguments to see what is mounted and where



```
user@cop4610: ~  
user@cop4610:~$ mount  
/dev/sda1 on / type ext3 (rw,errors=remount-ro)  
tmpfs on /lib/init/rw type tmpfs (rw,nosuid,mode=0755)  
proc on /proc type proc (rw,noexec,nosuid,nodev)  
sysfs on /sys type sysfs (rw,noexec,nosuid,nodev)  
procbususb on /proc/bus/usb type usbfs (rw)  
udev on /dev type tmpfs (rw,mode=0755)  
tmpfs on /dev/shm type tmpfs (rw,nosuid,nodev)  
devpts on /dev/pts type devpts (rw,noexec,nosuid,gid=5,mode=620)  
user@cop4610:~$
```

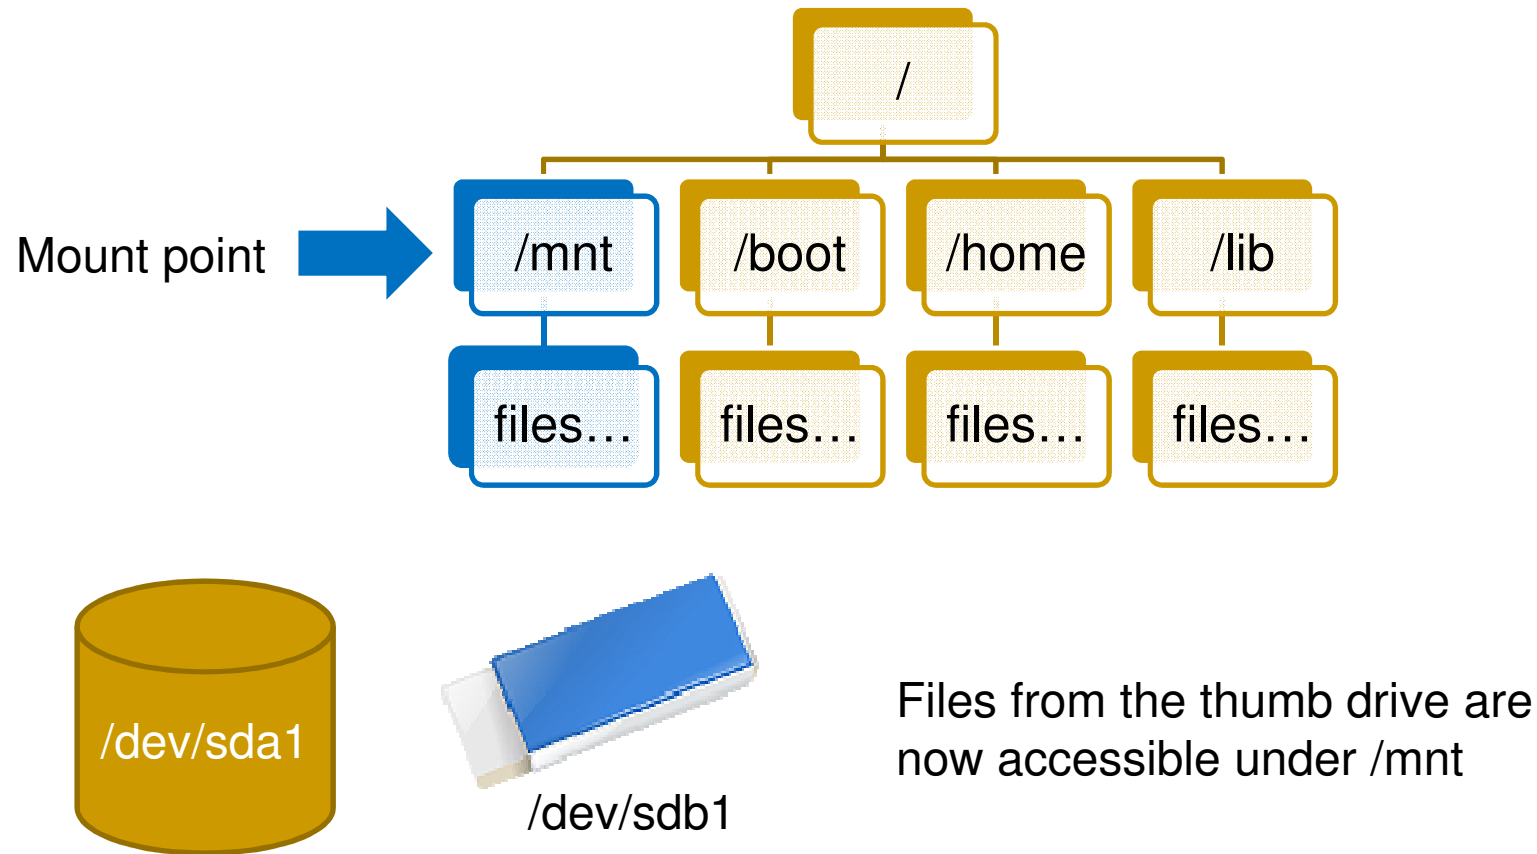
Root "/" file
system
mounted

Mount Example



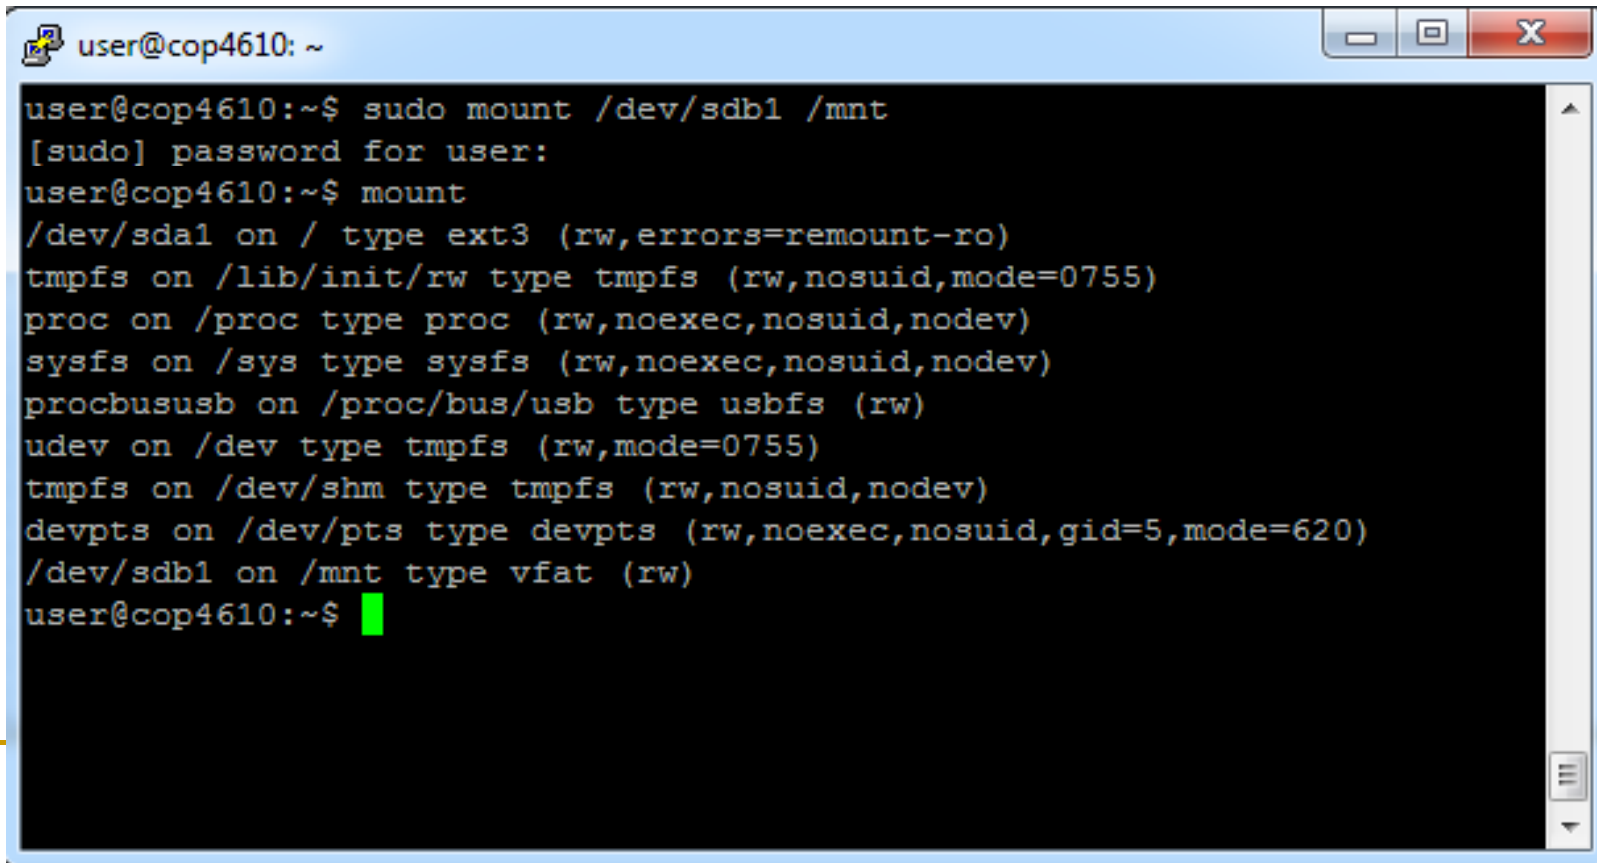
Now suppose we attach a thumb drive and want our thumb drive files accessible under `/mnt`...

File Hierarchy Example



Mount Example

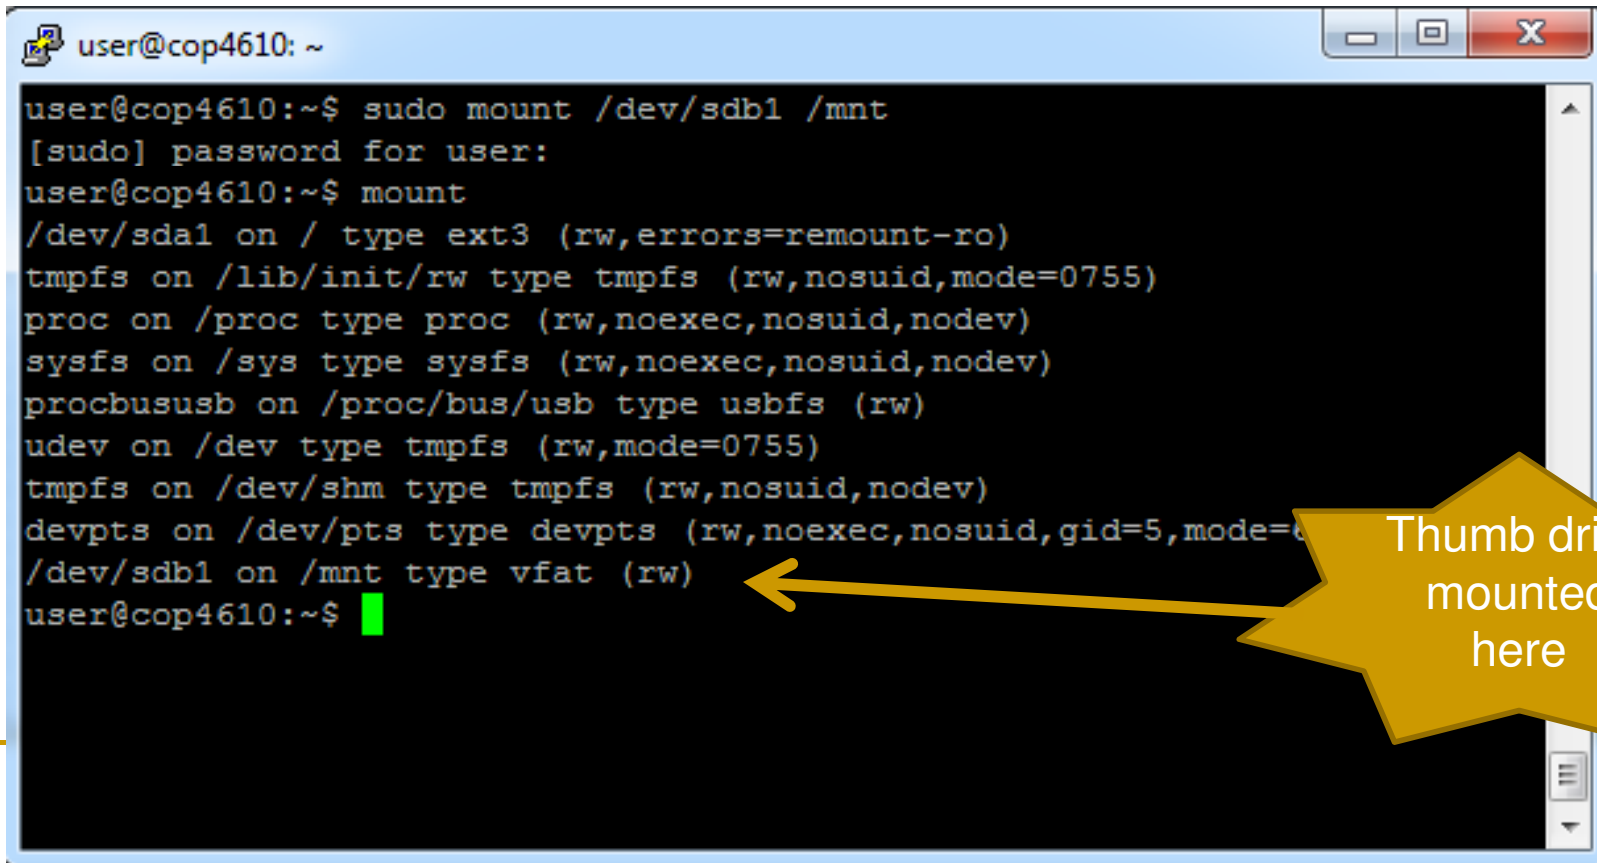
- The 'mount' command can dynamically attach new devices to new mount points

A terminal window titled 'user@cop4610: ~' with standard window controls. The terminal shows the execution of the 'mount' command. The user enters 'sudo mount /dev/sdb1 /mnt', followed by a password prompt. Then, the user enters 'mount', which displays a list of currently mounted filesystems, including the newly mounted '/dev/sdb1 on /mnt type vfat (rw)'.

```
user@cop4610:~$ sudo mount /dev/sdb1 /mnt
[sudo] password for user:
user@cop4610:~$ mount
/dev/sda1 on / type ext3 (rw,errors=remount-ro)
tmpfs on /lib/init/rw type tmpfs (rw,nosuid,mode=0755)
proc on /proc type proc (rw,noexec,nosuid,nodev)
sysfs on /sys type sysfs (rw,noexec,nosuid,nodev)
procbususb on /proc/bus/usb type usbfs (rw)
udev on /dev type tmpfs (rw,mode=0755)
tmpfs on /dev/shm type tmpfs (rw,nosuid,nodev)
devpts on /dev/pts type devpts (rw,noexec,nosuid,gid=5,mode=620)
/dev/sdb1 on /mnt type vfat (rw)
user@cop4610:~$
```

Mount Example

- The 'mount' command can dynamically attach new devices to new mount points



```
user@cop4610: ~  
user@cop4610:~$ sudo mount /dev/sdb1 /mnt  
[sudo] password for user:  
user@cop4610:~$ mount  
/dev/sda1 on / type ext3 (rw,errors=remount-ro)  
tmpfs on /lib/init/rw type tmpfs (rw,nosuid,mode=0755)  
proc on /proc type proc (rw,noexec,nosuid,nodev)  
sysfs on /sys type sysfs (rw,noexec,nosuid,nodev)  
procbususb on /proc/bus/usb type usbfs (rw)  
udev on /dev type tmpfs (rw,mode=0755)  
tmpfs on /dev/shm type tmpfs (rw,nosuid,nodev)  
devpts on /dev/pts type devpts (rw,noexec,nosuid,gid=5,mode=0660)  
/dev/sdb1 on /mnt type vfat (rw)  
user@cop4610:~$
```

Thumb drive mounted here

Un-mount Command

- `umount <dir>`

- In our example where the thumb drive was mounted at `/mnt`, we can issue

- `$> umount /mnt`

- Must have root privileges

Figuring out names of devices

- /etc/fstab – Has list of devices and file systems that get auto-mounted on boot
- 'dmesg' command shows output when plugging in a dynamic device

Project 3

More than you wanted to know about
FAT32..

Project 3

- You will create a user-space utility to manipulate a FAT32 file system image
 - No more kernel programming!
- Utility must understand a few basic commands to allow simple file system manipulation
- Utility must not corrupt the file system and should be robust

FAT32 Manipulation Utility

Utility only recognizes the following built-in commands:

- open
- close
- create
- rm
- size
- cd
- ls
- mkdir
- rmdir
- read
- write

File System Image

- Manipulation utility will work on a pre-configured FAT32 ***file system image***
 - Actually a file
- File system image will have raw FAT32 data structures inside
 - Just like looking at the raw bytes inside of a disk partition

File System Image

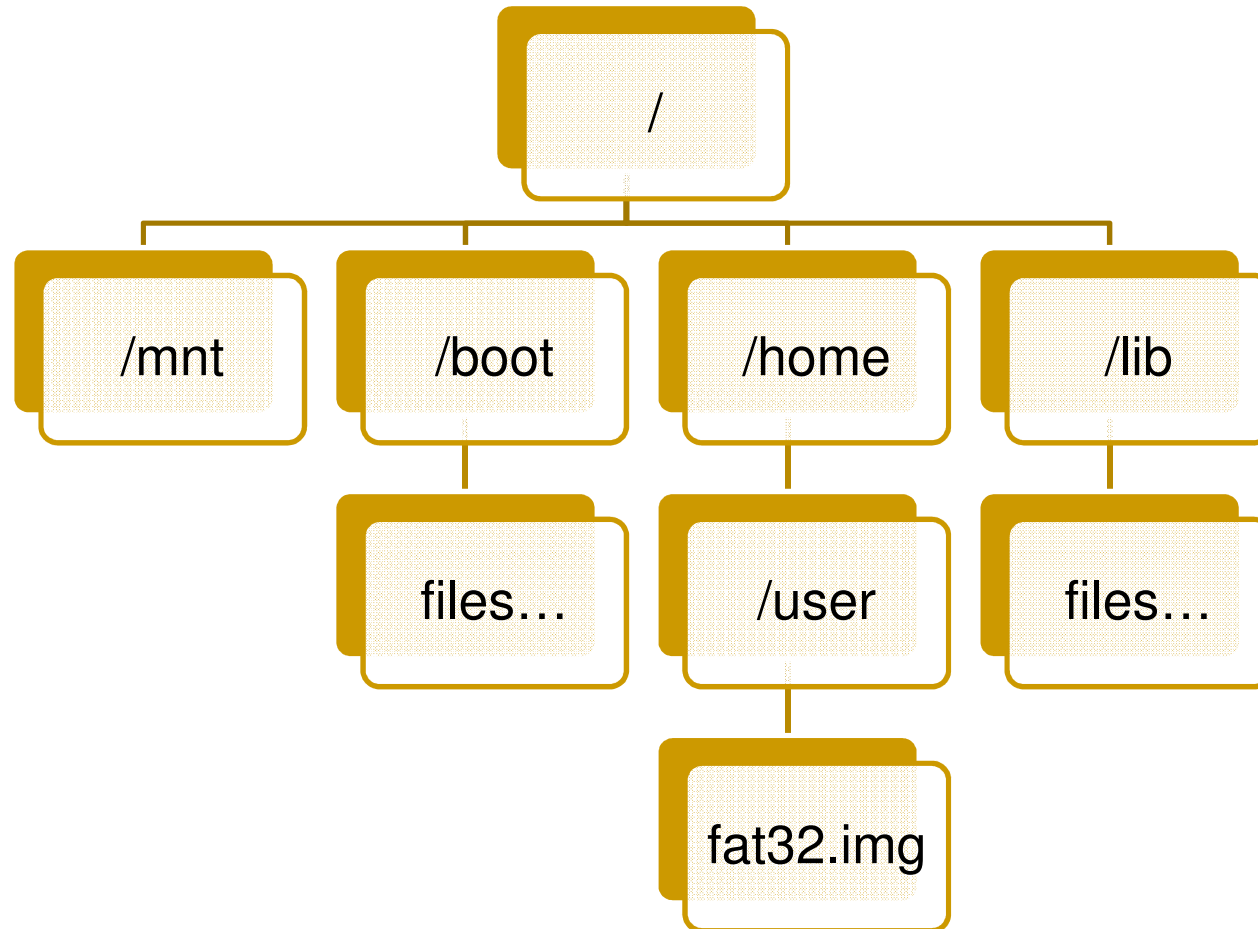
- Your FAT32 manipulation utility will have to
 - Open the FAT32 file system image
 - Read parts of the FAT32 file system image and interpret the raw bytes inside to service your utility's file system commands...

...just like a file system!

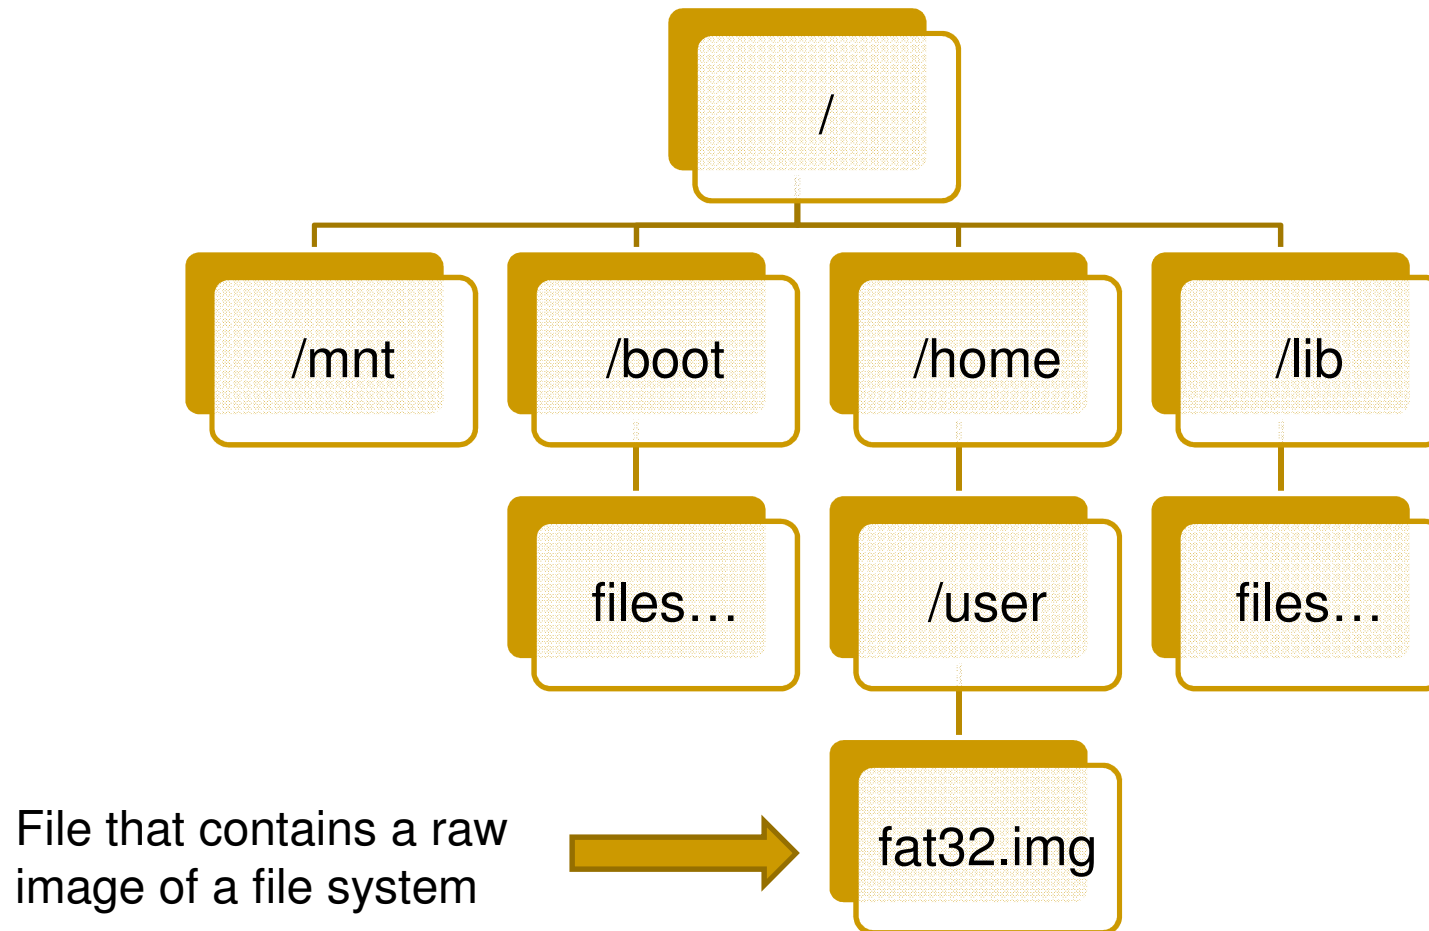
File System Image

- Sometimes you may want to check that you haven't corrupted your file system image, or that you can add or write files successfully
 - Mount your file system image with the OS FAT32 driver
 - Just like the file system image is a device

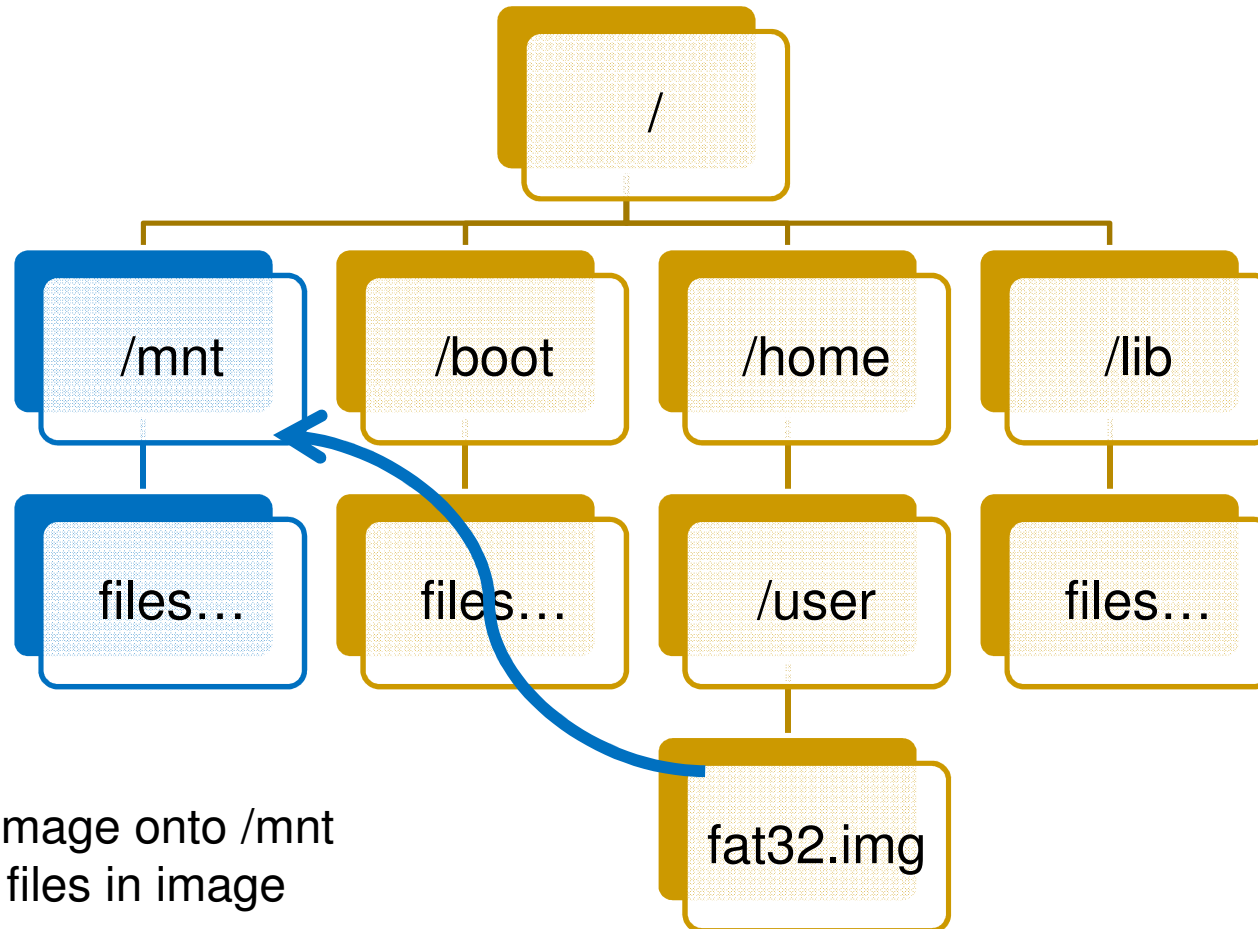
File System Image Mount Example



File System Image Mount Example



File System Image Mount Example



Mount image onto /mnt
to read files in image

File System Image Mount Example

```
$> sudo mount -o loop fat32.img /mnt  
$> cd /mnt
```

- `fat32.img` is your image file
- `/mnt` is your mounting directory
- Once the file is mounted, you can go into the `/mnt` directory and issue all your normal file system commands like:
 - `ls`, `cat`, `cd`, ...


General FAT32 Data Structures

Terminology

- **Byte** – 8 bits of data, the smallest addressable unit in modern processors
- **Sector** – Smallest addressable unit on a storage device. Usually this is 512 bytes
- **Cluster** – FAT32-specific term. A group of sectors representing a chunk of data
- **FAT** – Stands for *file allocation table* and is a map of files to data

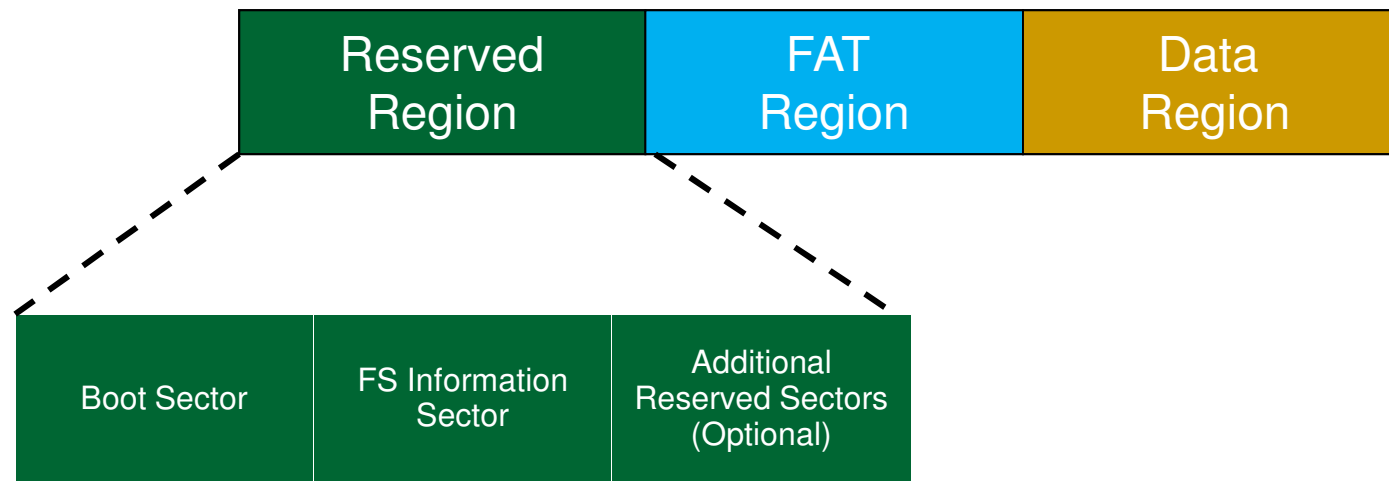
FAT32 Disk Layout

- 3 main regions...



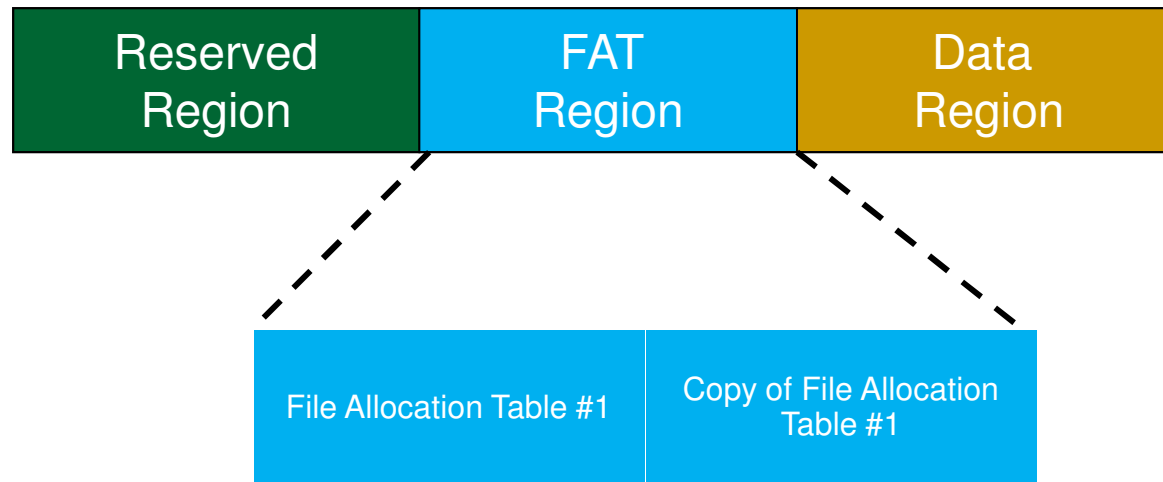
Reserved Region

- **Reserved Region** – Includes the boot sector, the extended boot sector, the file system information sector, and a few other reserved sectors




FAT Region

- **FAT Region** – A map used to traverse the data region. Contains mappings from cluster locations to cluster locations



Data Region

- **Data Region** – Using the addresses from the FAT region, contains actual file/directory data



Endian

Big or little?

Machine Endianness

- The endianness of a given machine determines in what order a group of bytes are handled (ints, shorts, long longs)
 - Big-endian – most significant byte first
 - Little-endian – least significant byte first
 - This is important to understand for this project, since FAT32 is always formatted as little-endian
-

FAT32 Endianness

- The following are a few cases where endianness matters in your project:
 - ❑ Reading in integral values from the FAT32 image
 - ❑ Reading in shorts from a FAT32 image
 - ❑ Combining multiple shorts to form a single integer from the FAT32 image
 - ❑ Interpreting directory entry attributes
-

Endian Example (English Version)

- Imagine you can only communicate three letters at a time, and your word is “RAPID”
 - Big-endian
 - ❑ 1. RAP
 - ❑ 2. ID
 - ❑ Word = RAPID
 - Little-endian
 - ❑ 1. PID
 - ❑ 2. RA
 - ❑ Word = PIDRA (come again?)
-

Endian Example (data version)

- `short value = 15; /* 0x000F */`
- `char bytes[2];`
- `memcpy(bytes, &value, sizeof(short));`
- In little-endian:
 - `bytes[0] = 0x0F`
 - `bytes[1] = 0x00`
- In big-endian:
 - `bytes[0] = 0x00`
 - `bytes[1] = 0x0F`

Endian Example (data version 2)

- `int value = 13371337; /* 0x00CC07C9 */`
- `char bytes[4];`
- `memcpy(bytes, &value, sizeof(int));`
- In little-endian:
 - ❑ `bytes[0] = 0xC9`
 - ❑ `bytes[1] = 0x07`
 - ❑ `bytes[2] = 0xCC`
 - ❑ `bytes[3] = 0x00`
- In big-endian:
 - ❑ `bytes[0] = 0x00`
 - ❑ `bytes[1] = 0xCC`
 - ❑ `bytes[2] = 0x07`
 - ❑ `bytes[3] = 0x09`

Visualizing Example 2

Value = 13371337 (0x00CC07C9)

index	0	1	2	3
little endian	0xC9	0x07	0xCC	0x00
big endian	0x00	0xCC	0x07	0xC9

Additional Project 3 Information

- Like other projects, may work in teams or alone
- Project deadline is December 3rd, 2010 at 11:59:59pm (Friday before finals week)
 - No extensions
 - Can use remainder of slack days
 - Cannot be more than 3 days late
 - Monday, December 6th is absolute last day to turn in

Project 3 Partners

- Raise your hand if you are looking for a partner...
- Also use the discussion board

Until Next Time

- Set up your environment
- Download the image file
- Practice mounting the image file with the OS FAT32 drivers
 - Make sure you can cd into /mnt and read/write to the files
- Read over the FAT32 Specification